

Simulated Annealing

Trying to find good solutions from bad beginnings....



A sample problem, first

- Let's say that you've got a function, and you're trying to find a "best answer" from a set of inputs
 - For instance, trying to solve the Knapsack problem with a maximum weight of 1000 units, and 500 items of varying weights to choose from
 - The "traveling salesman" problem is another example: given a number of cities and the costs of traveling from any city to any other city, what is the cheapest round-trip route that visits each city exactly once and then returns to the starting city?
- How do you find the best answer?



Some possible approaches

- You can try combinations of stuff at random
 - "I'm going to try 2000 combinations, and just use the best one"
 - Doesn't guarantee a good solution at all (might be way too light, might go over on all of them, etc.)
- You can "hill climb"
 - Sort the objects based on size ($O(n \log n)$) and then keep putting the "next heaviest thing that will fit" into the knapsack ($O(n)$)
 - Should give you some sort of an approximation, but no guarantee of a perfect fit
 - A decent, but "greedy" method
- You can try all possible combinations
 - Guaranteed to work, but explosive complexity ($O(n!)$)
-

Or, you can try something different

Enter simulated annealing (SA)


- Simulated annealing (SA) is one approach that can be used when trying to find a good approximation to a global optimum of a given function in a large search space.
- In other words:
 - If you're looking for a minimum (or maximum) value from a function given a lot of possible inputs, SA is one way of trying to get a good approximation.

Where does it come from?

- Annealing:
 - A process in metallurgy where a metal is first heated up and then put through a controlled cooling process
 - The idea is that this increases the internal crystalline structure of the metal, making it harder/stronger, and decreasing the number of defects
 - The heat causes the atoms to become "unstuck" from their initial positions (a local minimum of the internal energy) and wander randomly through states of higher energy
 - The slow cooling gives them more chances of finding configurations with lower internal energy than the initial one

- In simulated annealing:
 - We generate an initial configuration for the inputs to the function
 - We then pick a new configuration at random from "nearby" the current one, and see if it yields a better result
 - If it does, we'll (usually) make this our new "current configuration", and continue on
 - If it *doesn't* give us a better result, we might still make it our current configuration, but we make that choice at random
 - The "heating" and "cooling" processes are what determine whether or not we'll use a poorer solution as a current configuration
 - Early on (while the "temperature" is still high), we are more likely to allow a poorer solution to be used
 - As the algorithm continues (and the "temperature" falls), we become less and less likely to accept a poorer solution.

- Using SA means that we need to be able to:
 - Generate "neighbors" of a given configuration quickly/easily
 - In Knapsack, this might be done by adding/removing a single item from the collection being considered
 - In Traveling Salesman, we might pick two of the cities in the current travel plan and swap the order in which they are visited
 - Calculate the "score" for a given configuration
 - In Knapsack, this might be done by calculating how much more stuff can be put into the sack
 - In Traveling Salesman, this could simply be the length of the current travel plan
 - Traditionally, SA calls this the "energy" of the configuration, and tries to find the "lowest energy" solution



Allowing for "downhill" moves saves the us from getting stuck at local maxima, which can happen with a "greedy" method.

The core algorithm

```
// Figure out where we're starting from
curConfig = generateRandomConfig();
curScore = score(curConfig)           // Remember, lower scores are better...

// It's the best that we've seen so far...
bestConfig = curConfig
bestScore = curScore

k = 0 // Energy evaluation count
while k < kmax and curScore > maxAcceptableScore
    newConfig = randomNeighbour(config) // Pick some neighbour
    newScore = score(newConfig)         // Compute its energy
    if (newScore < bestScore) {
        // Save the "new best" configuration's data
        bestConfig = newConfig
        bestScore = newScore
    }
    if (random() < calculateProbability(curScore, newScore, ((float)k)/kmax)) {
        // Make it the new configuration
        curConfig = newConfig
        curScore = newScore
    }
    k = k + 1 // One more evaluation done
}


// Return the best solution found
return bestConfig
```

An example of calculateProbability()

```
float calculateProbability( curScore, newScore, temperature ) {
    if ( newScore < curScore ) {
        // lower scores are better, so the new configuration is an improvement!
        return 1.0
    }
    else {
        // If we're worse off with the new config, we're less likely to accept it
        // The "colder" it gets, the less likely we will accept it
        return e^(curScore - newScore)/temperature);
    }
}
```

Some extensions to SA

- "Restarting"
 - Sometimes we might decide that we're just too far off in the weeds (e.g., we've accepted a series of poorer solutions)
 - In this case, we might move back to the best configuration seen so far, and maybe raise the temperature somewhat (to increase randomness again)
- "Tabu Search" (TA)
 - In simple SA, it's possible for the system to get caught in a cycle, where it keeps moving between a group of possible configurations
 - In Tabu Search, we keep a "tabu list" of configurations that we've already seen (e.g., the last 10), and moves to those configurations are not allowed



Any questions?
