

File Organization

Reminder

- Final exam
 - The date for the Final has been decided:
 - Saturday, November 16th
 - 8am – 10am
 - 01-2000

Project Notes

- Change your generic solver?
 - Don't forget to change your Clock and Farmer problem as well.
- Memory Management
 - Using purify
 - Add to Makefile:
 - CCC = purify CC
 - Or better yet, create a file header.mak
 - Workshop also has memory management tools.
- Parking Lot Problem: due Nov 11th

New plan

- Today: Files 1
- Tuesday: Exam / Files 2
- Thursday: Files 3

- Monday: Ethics
- Tuesday: Final Review
- Then we are done!

Before we start


- Any questions

File Organization

- How to find stuff in a file
 - Improve access time by imposing a defined structure on a file
 - Database applications
 - Advanced searching strategies

It is always best to start at the beginning...

- Hard Disk – Physical layout



Hard Disk – Physical layout

- Hard disk
 - Is composed of a number of cylinder
 - Each cylinder is composed of tracks
 - Each track is composed of sectors
 - Sector is the lowest unit of read/writable data
 - Each sector can be addressed by giving
 - Cylinder number
 - Track number
 - Sector number
 - Info on a sector is called a block of data

Hard Disk – Performance


- Seek Time
 - time to reposition the head and increases with the distance that the head must move.
- Rotational latency time
 - time from when the head is over the correct track until the data rotates around and is under the head and can be read.
- Access Time
 - time from when a read or write request is issued to when the data transfer begins. It is the sum of the seek time and latency time.

Hard Disk – Performance

- Data Transfer Rate
 - rate at which data can be retrieved from the disk and sent to the controller.
- Mean Time to failure
 - is the number of hours (on average) until a disk fails.

Going up a level

- It is the responsibility of the Operating System to:
 - Map logical files to physical sector I/O.
 - Return blocks of data to a program.



Going up a level

- The OS must also:
 - Organize it's file system for maximum performance (I.e. low access time)
 - File System
- If the file is a searchable database
 - This optimization can be improved by imposing an additional structure on a file.

Terminology

- Record
 - Collection of data pertaining to one entity
 - E.g. Employee
 - Each record consists of a number of data fields.
 - Files are composed of a collection of records.
- Key
 - Field within a record upon which a search is made

Organization Strategies

- Sequential
- Indexed
- Hashing
- Tree-based Organization

Sequential Organization

- Simplest organization
 - Place records one right after another
 - Array / linked list of records

Sequential Organization

- Can order info by sorting records by key.
 - Search for record with given key is $O(n)$
 - 1,000,000 records
 - on avg 500,000 records reads to find desired record
 - Can this be improved by doing a binary search?

Only if underlying OS allows for random access

Sequential Organization

- Problem with sequential organization
 - Search is $O(n)$.
 - Increases linearly as number of records increase.


Indexed Files

- Construct an index
 - The index contains:
 - Each record
 - Record's location in the file.
 - Dense index – contains every key for every record
 - Searching the file involves
 - Searching the index for the key
 - Directly accessing the record based on the record's address in the index file.

Indexed Files

- To be efficient:
 - Index must be sorted by key
 - Index file must support direct access
 - Data file must support direct access
- Note that data need not be sorted, just the index.
- Searching the index is more efficient than searching the file since the index file will be smaller than the file being indexed.

Indexed Files



Indexed Files

- Example:
 - 1,000,000 records
 - 10 key/location pairs fit in a block
 - Index file 100,000 blocks long
 - On avg, search index would require 50,000 index block reads + 1 data block read.


Indexed Files

- Problem with indexed files
 - Index contains key for each record
 - As number of records grow, so does the size of the index file.
 - Index file can become unmanageably large.

Indexed Sequential Files

- Variation on the indexed file
 - Index file is sparse
 - Data is stored sequentially by key
- Index files indexes groups of records
 - Index gives location of a group of records
 - Index can point to first record in a group
 - Index can point to last record in a group
 - Sequential search within each group

Indexed Sequential Files



Indexed Sequential Files

- Example
 - 1,000,000 records
 - 10 key/location pairs fit in a block
 - Every 10th record key is placed in index
 - Index size: 10,000 records
 - On average: 5,000 index block reads + 5 data record reads

Indexed Sequential Files

- For efficiency
 - Sparse index can be loaded into memory (if it is small enough).


Searching Indexed Files

- Dense Index:
 - If key is found in index, location of record is given
 - Else the desired record is not present
- Sparse Index
 - If key is in index, location of record is given
 - Else index indicates group where record should be found
 - Sequential search from start of record group to end of group.
- Questions?

Multi-Level Indices

- If data file becomes unmanageably large:
 - Construct a sparse index of the data
 - Construct a sparse index of the sparse index.

Multi-Level Indices



Multi-Level Indices

- Searching
 - Search primary index for where to start searching the secondary index
 - Search secondary index for where to start searching data.
 - Search data sequentially from location given in secondary index.
- Use as many levels as needed

Multi-Level Indices

- Example
 - 1,000,000 records
 - 10 key/location pairs per record
 - Every 10th key in index
 - Create 3 index files:
 - 100,000 entries
 - 10,000 entries
 - 1,000 entries
 - On average: 500 blocks from 1st index, 5 from second, 5 from third, 5 data records.

Multi-Level Indices

- Commercial Implementation
 - ISAM (Indexed Sequential Access Method) from IBM
 - Records blocked into tracks
 - Tracks grouped into cylinders
 - Each disk has a cylinder index
 - If multiple disks, there was a disk index
 - Useful for “static” files
 - Additions / deletions expensive
- Questions?

Sequential and Indexed Files


Organization	Records / avg search
Sequential	500,000
Indexed	50,001
Indexed Sequential	5,005
Multi-Level Index	515

Sequential and Indexed Files

- Questions?

Hashing

- Recall from CS2



Hashing

- Hashing Schemes
 - Each bucket corresponds to a sector
 - Best if 1 record fits nicely in a block
 - For large records (n blocks / record)
 - Hash to first block of record
 - For small record (< 1 block/ record)
 - Hash to record
 - Use extra space for collision management

Summary

- File organization Schemes
 - For efficient access to data records
- Organizations
 - Sequential
 - Indexed
 - Dense index
 - Indexed Sequential (Sparse Index)
 - Multi-level
 - Hashing
 - Next time: Tree-based organizations