

Abstract Data Types and Data Structures

Computer Science S-111
Harvard University


David G. Sullivan, Ph.D.

Congrats on completing the first half!

- In the second half, we will study fundamental *data structures*.
 - ways of imposing order on a collection of information
 - sequences: lists, stacks, and queues
 - trees
 - hash tables
 - graphs
- We will also:
 - study *algorithms* related to these data structures
 - learn how to *compare* data structures & algorithms
- Goals:
 - learn to think more intelligently about programming problems
 - acquire a set of useful tools and techniques

Sample Problem I: Finding Shortest Paths

- Given a set of routes between pairs of cities, determine the shortest path from city A to city B.



Sample Problem II: A Data "Dictionary"

- Given a large collection of data, how can we arrange it so that we can efficiently:
 - add a new item
 - search for an existing item
- Some data structures provide better performance than others for this application.
- More generally, we'll learn how to characterize the *efficiency* of different data structures and their associated algorithms.

Example of Comparing Algorithms

- Consider the problem of finding a phone number in a phonebook.
- Let's informally compare the time efficiency of two algorithms for this problem.

Algorithm 1 for Finding a Phone Number

```
findNumber(person) {  
    for (p = number of first page; p <= number of the last page; p++) {  
        if person is found on page p {  
            return the person's phone number  
        }  
    }  
    return NOT_FOUND  
}
```

- If there were 1,000 pages in the phonebook, how many pages would this look at in the worst case?
- What if there were 1,000,000 pages?

Algorithm 2 for Finding a Phone Number

```
findNumber(person) {  
  min = the number of the first page  
  max = the number of the last page  
  while (min <= max) {  
    mid = (min + max) / 2    // page number of the middle page  
    if person is found on page mid {  
      return the person's number  
    } else if the person's name comes earlier in the book {  
      max = mid - 1  
    } else {  
      min = mid + 1  
    }  
  }  
  return NOT_FOUND  
}
```

- If there were 1,000 pages in the phonebook, how many pages would this look at in the worst case?
- What if there were 1,000,000 pages?

Searching a Collection of Data

- The phonebook problem is one example of a common task: searching for an item in a collection of data.
 - another example: searching for a record in a database
- Algorithm 1 is known as *sequential search*.
 - also called *linear search*
- Algorithm 2 is known as *binary search*.
 - only works if the items in the data collection are sorted

Abstract Data Types

- An *abstract data type* (ADT) is a model of a data structure that specifies:
 - the characteristics of the collection of data
 - the operations that can be performed on the collection
- It's *abstract* because it doesn't specify *how* the ADT will be implemented.
 - does *not* commit to any low-level details
- A given ADT can have multiple implementations.

A Simple ADT: A Bag

- A bag is just a container for a group of data items.
 - analogy: a bag of candy
- The positions of the data items don't matter (unlike a list).
 - {3, 2, 10, 6} is equivalent to {2, 3, 6, 10}
- The items do *not* need to be unique (unlike a set).
 - {7, 2, 10, 7, 5} isn't a set, but it is a bag

A Simple ADT: A Bag (cont.)

- The operations we want a Bag to support:
 - `add(item)`: add `item` to the Bag
 - `remove(item)`: remove one occurrence of `item` (if any) from the Bag
 - `contains(item)`: check if `item` is in the Bag
 - `numItems()`: get the number of items in the Bag
 - `grab()`: get an item at random, without removing it
 - reflects the fact that the items don't have a position (and thus we can't say "get the 5th item in the Bag")
 - `toArray()`: get an array containing the current contents of the bag
- We want the bag to be able to store objects of any type.

Specifying an ADT Using an Interface

- In Java, we can use an *interface* to specify an ADT:

```
public interface Bag {
    boolean add(Object item);
    boolean remove(Object item);
    boolean contains(Object item);
    int numItems();
    Object grab();
    Object[] toArray();
}
```
- An interface specifies a set of methods.
 - includes only the method headers
 - does *not* typically include the full method definitions
- Like a class, it must go in a file with an appropriate name.
 - in this case: `Bag.java`

Implementing an ADT Using a Class

- To implement an ADT, we define a class:

```
public class ArrayBag implements Bag {  
    ...  
    public boolean add(Object item) {  
        ...  
    }  
}
```

- When a class header includes an `implements` clause, the class must define all of the methods in the interface.
 - if the class doesn't define them, it won't compile

All Interface Methods Are Public

- Methods specified in an interface *must* be public, so we don't use the keyword `public` in the definition:

```
public interface Bag {  
    boolean add(Object item);  
    boolean remove(Object item);  
    boolean contains(Object item);  
    int numItems();  
    Object grab();  
    Object[] toArray();  
}
```

- However, when we actually implement the methods in a class, we *do* need to use `public`:

```
public class ArrayBag implements Bag {  
    ...  
    public boolean add(Object item) {  
        ...  
    }  
}
```

One Possible Bag Implementation

- One way to store the items in the bag is to use an array:

```
public class ArrayBag {  
    private _____[] items;  
  
    ...  
}
```

- What type should the array be?
- This allows us to store *any* type of object in the `items` array, thanks to the power of polymorphism:

```
ArrayBag bag = new ArrayBag();  
bag.add("hello");  
bag.add(new Rectangle(20, 30));
```

- How could we keep track of how many items are in a bag?

Another Example of Polymorphism

- An interface name can be used as the type of a variable:

```
Bag b;
```

- Variables with an interface type can refer to objects of any class that implements the interface:

```
Bag b = new ArrayBag();
```

- Using the interface as the type allows us to write code that works with any implementation of an ADT:

```
public void processBag(Bag b) {  
    for (int i = 0; i < b.numItems(); i++) {  
        ...  
    }  
}
```

- the param can be an instance of *any* Bag implementation
- we must use method calls to access the object's internals, because the fields are not part of the interface

Memory Management: Looking Under the Hood

- To understand how data structures are implemented, you need to understand how memory is managed.
- There are three main types of memory allocation in Java.
- They correspond to three different regions of memory.

Memory Management, Type I: Static Storage

- Static storage is used for *class variables*, which are declared *outside any method* using the keyword `static`:


```
public class MyMethods {  
    public static int numCompares;  
    public static final double PI = 3.14159;
```

- There is only one copy of each class variable.
 - shared by all objects of the class
 - Java's version of a global variable
- The Java runtime allocates memory for class variables when the class is first encountered.
 - this memory stays fixed for the duration of the program

Memory Management, Type II: Stack Storage

- Method parameters and local variables are stored in a region of memory known as *the stack*.
- For each method call, a new *stack frame* is added to the top of the stack.

```
public class Foo {  
    public static int x(int i) {  
        int j = i - 2;  
        if (i >= 6) {  
            return i;  
        }  
        return x(i + j);  
    }  
    public static void  
    main(String[] args) {  
        System.out.println(x(5));  
    }  
}
```



- When a method completes, its stack frame is removed.

Memory Management, Type III: Heap Storage

- Objects are stored in a memory region known as *the heap*.
- Memory on the heap is allocated using the `new` operator:

```
int[] values = new int[3];  
ArrayBag b = new ArrayBag();
```

- `new` returns the memory address of the start of the object on the heap.
 - a reference!
- An object stays on the heap until there are no remaining references to it.
- Unused objects are automatically reclaimed by a process known as *garbage collection*.
 - makes their memory available for other objects

Two Constructors for the ArrayBag Class

```
public class ArrayBag implements Bag {
    private Object[] items;
    private int numItems;
    public static final int DEFAULT_MAX_SIZE = 50;

    public ArrayBag() {
        this.items = new Object[DEFAULT_MAX_SIZE];
        this.numItems = 0;
    }
    public ArrayBag(int maxSize) {
        ...
    }
}
```

- As we've seen before, we can have multiple constructors.
 - the parameters must differ in some way
- The first one is useful for small bags.
 - creates an array with room for 50 items.
- The second one allows the client to specify the max # of items.

Two Constructors for the ArrayBag Class

```
public class ArrayBag implements Bag {
    private Object[] items;
    private int numItems;
    public static final int DEFAULT_MAX_SIZE = 50;

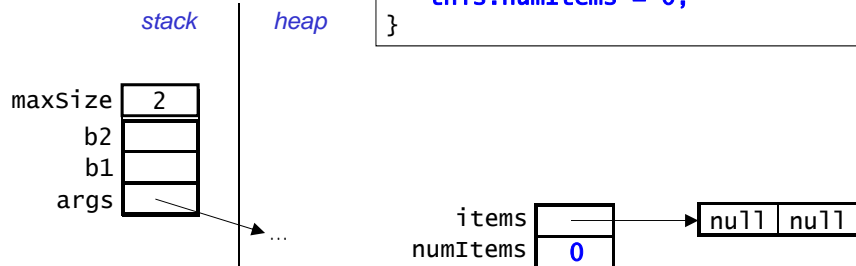
    public ArrayBag() {
        this.items = new Object[DEFAULT_MAX_SIZE];
        this.numItems = 0;
    }
    public ArrayBag(int maxSize) {
        if (maxSize <= 0) {
            throw new IllegalArgumentException(
                "maxSize must be > 0");
        }
        this.items = new Object[maxSize];
        this.numItems = 0;
    }
    ...
}
```

- If the user inputs an invalid maxSize, we throw an exception.

Example: Creating Two ArrayBag Objects

```
// client
public static void main(String[] args) {
    ArrayBag b1 = new ArrayBag(2);
    ArrayBag b2 = new ArrayBag(4);
    ...
}
```

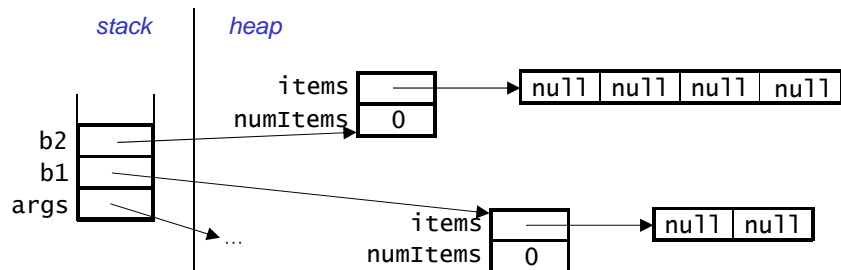
```
// constructor
public ArrayBag(int maxSize) {
    ... // error-checking
    this.items = new Object[maxSize];
    this.numItems = 0;
}
```



Example: Creating Two ArrayBag Objects


```
// client
public static void main(String[] args) {
    ArrayBag b1 = new ArrayBag(2);
    ArrayBag b2 = new ArrayBag(4);
    ...
}
```

- After the objects have been created, here's what we have:




Adding Items


- We fill the array from left to right. Here's an empty bag:



- After adding the first item:




- After adding the second item:




Adding Items (cont.)

- After adding the third item:



- After adding the fourth item:



- At this point, the ArrayBag is full!
 - it's non-trivial to "grow" an array, so we don't!
 - additional items cannot be added until one is removed

A Method for Adding an Item to a Bag

```
public class ArrayBag implements Bag {
    private Object[] items;
    private int numItems;
    ...
    public boolean add(Object item) {
        if (item == null) {
            throw new IllegalArgumentException("no nulls");
        } else if (this.numItems == this.items.length) {
            return false; // no more room!
        } else {
            this.items[this.numItems] = item;
            this.numItems++;
            return true; // success!
        }
    }
    ...
}
```

The diagram shows a variable `items` pointing to an array of four slots. The first slot contains the string "hello, world", the second contains "howdy", the third contains "bye", and the fourth contains "see ya!". Below the array, the variable `numItems` is shown with the value 4, indicating that all slots in the array are currently occupied.

A Method for Adding an Item to a Bag


```
public class ArrayBag implements Bag {
    private Object[] items;
    private int numItems;
    ...
    public boolean add(Object item) {
        if (item == null) {
            throw new IllegalArgumentException("no nulls");
        } else if (this.numItems == this.items.length) {
            return false; // no more room!
        } else {
            this.items[this.numItems] = item;
            this.numItems++;
            return true; // success!
        }
    }
    ...
}
```

- Initially, `this.numItems` is 0, so the first item goes in position 0.
- We increase `this.numItems` because we now have 1 more item.
 - and so the *next* item added will go in the correct position!

Example: Adding an Item

```
public static void main(String[] args) {
    String message = "hello, world";
    ArrayBag b = new ArrayBag(4);
    b.add(message);
    ...
}
```


```
public boolean add(Object item) {
    ...
    else {
        this.items[this.numItems] = item;
        this.numItems++;
        return true;
    } ...
}
```



Example: Adding an Item (cont.)

```
public static void main(String[] args) {
    String message = "hello, world";
    ArrayBag b = new ArrayBag(4);
    b.add(message);
    ...
}
```

```
public boolean add(Object item) {
    ...
    else {
        this.items[this.numItems] = item;
        this.numItems++;
        return true;
    } ...
}
```




- `add`'s stack frame includes:
 - `item`, which stores...
 - `this`, which stores...

Example: Adding an Item (cont.)

```
public static void main(String[] args) {
    String message = "hello, world";
    ArrayBag b = new ArrayBag(4);
    b.add(message);
    ...
}
```

```
public boolean add(Object item) {
    ...
    else {
        this.items[this.numItems] = item;
        this.numItems++;
        return true;
    } ...
}
```




- The method modifies the `items` array and `numItems`.
 - note that the array holds a copy of the *reference* to the item, not a copy of the item itself.

Example: Adding an Item (cont.)


```
public static void main(String[] args) {
    String message = "hello, world";
    ArrayBag b = new ArrayBag(4);
    b.add(message);
    ...
}
```

```
public boolean add(Object item) {
    ...
    else {
        this.items[this.numItems] = item;
        this.numItems++;
        return true;
    } ...
}
```



- After the method call returns, `add`'s stack frame is removed from the stack.

Extra Practice: Determining if a Bag Contains an Item




- Let's write the `ArrayBag` `contains()` method together.
 - should return `true` if an object equal to `item` is found, and `false` otherwise.

```
_____ contains(_____ item) {
```

```
}
```

Would this work instead?



- Let's write the `ArrayBag` `contains()` method together.
 - should return `true` if an object equal to `item` is found, and `false` otherwise.

```
public boolean contains(Object item) {  
    for (int i = 0; i < this.items.length; i++) {  
        if (this.items[i].equals(item)) { // not ==  
            return true;  
        }  
    }  
    return false;  
}
```

Another Incorrect `contains()` Method

```
public boolean contains(Object item) {  
    for (int i = 0; i < this.numItems; i++) {  
        if (this.items[i].equals(item)) {  
            return true;  
        } else {  
            return false;  
        }  
    }  
    return false;  
}
```

- What's the problem with this?

A Method That Takes a Bag as a Parameter

```
public boolean containsAll(Bag otherBag) {
    if (otherBag == null || otherBag.numItems() == 0) {
        return false;
    }
    Object[] otherItems = otherBag.toArray();
    for (int i = 0; i < otherItems.length; i++) {
        if (! this.contains(otherItems[i])) {
            return false;
        }
    }
    return true;
}
```

- We use Bag instead of ArrayBag as the type of the parameter.
 - allows this method to be part of the Bag interface
 - allows us to pass in *any* object that implements Bag
- We must use methods in the interface to manipulate otherBag.
 - we can't use the fields, because they're not in the interface

A Type Mismatch

- Here are the headers of two ArrayBag methods:

```
public boolean add(Object item)
public Object grab()
```
- Polymorphism allows us to pass String objects into add():

```
ArrayBag stringBag = new ArrayBag();
stringBag.add("hello");
stringBag.add("world");
```
- However, this will not work:

```
String str = stringBag.grab(); // compiler error
```

 - the return type of grab() is Object
 - Object isn't a subclass of String, so polymorphism doesn't help!
- Instead, we need to use a *type cast*:

```
String str = (String)stringBag.grab();
```

 - this cast doesn't actually change the value being assigned
 - it just reassures the compiler that the assignment is okay